

# 2-linear-regression-pytorch

December 25, 2021

## 1 Implementation

```
[15]: # Import Numpy & PyTorch
import numpy as np
import torch
```

### 1.1 Linear Regression Model using PyTorch built-ins

Let's re-implement the same model using some built-in functions and classes from PyTorch.

And now using two different targets: Apples and Oranges

```
[16]: # Imports
import torch.nn as nn
```

```
[17]: # Input (temp, rainfall, humidity)
inputs = np.array([[73, 67, 43], [91, 88, 64], [87, 134, 58], [102, 43, 37],
→[69, 96, 70], [73, 67, 43], [91, 88, 64], [87, 134, 58], [102, 43, 37], [69,
→96, 70], [73, 67, 43], [91, 88, 64], [87, 134, 58], [102, 43, 37], [69, 96,
→70]], dtype='float32')
# Targets (apples, oranges)
targets = np.array([[56, 70], [81, 101], [119, 133], [22, 37], [103, 119],
→[56, 70], [81, 101], [119, 133], [22, 37], [103, 119],
→[56, 70], [81, 101], [119, 133], [22, 37], [103, 119]],
→dtype='float32')
```

```
[18]: inputs = torch.from_numpy(inputs)
targets = torch.from_numpy(targets)
```

#### 1.1.1 Dataset and DataLoader

We'll create a TensorDataset, which allows access to rows from inputs and targets as tuples. We'll also create a DataLoader, to split the data into batches while training. It also provides other utilities like shuffling and sampling.

```
[19]: # Import tensor dataset & data loader
from torch.utils.data import TensorDataset, DataLoader
```

```
[20]: # Define dataset
```

```
[21]: # Define data loader
```

### 1.1.2 nn.Linear

Instead of initializing the weights & biases manually, we can define the model using `nn.Linear`.

```
[22]: # Define model
```

### 1.1.3 Optimizer

Instead of manually manipulating the weights & biases using gradients, we can use the optimizer `optim.SGD`.

```
[ ]: #Obtain Parameters
```

```
[23]: # Define optimizer
```

### 1.1.4 Loss Function

Instead of defining a loss function manually, we can use the built-in loss function `mse_loss`.

```
[24]: # Import nn.functional
import torch.nn.functional as F
```

```
[25]: # Define loss function
loss_fn = F.mse_loss
```

```
[26]: #loss = loss_fn(?, ?)
#print(loss)
```

### 1.1.5 Train the model

We are ready to train the model now. We can define a utility function `fit` which trains the model for a given number of epochs.

```
[27]: # Define a utility function to train the model
#def fit(num_epochs, model, loss_fn, opt):
#    for epoch in range(num_epochs):
#        for xb, yb in train_dl:
#            # Generate predictions
#            pred = model(?)
#            loss = loss_fn(?, ?)
#            # Perform gradient descent
#            loss.backward()
#            opt.step()
#            opt.zero_grad()
#        print('Training loss: ', loss_fn(model(inputs), targets))
```

```
[28]: # Train the model for 100 epochs
      #fit(?, ?, ?, ?)
```

```
[29]: # Generate predictions
      #preds = model(?)
      #preds
```

```
[30]: # Compare with targets
      #targets
```

Now we can define the model, optimizer and loss function exactly as before.

```
[31]: #fit(?, ?, ?, ?)
```

#Exercise 1: Try Linear Regression just using numpy (Without Tensorflow/Pytorch or other torch library). You can optionally use sklearn (if you want)

Exercise 2: Try Linear regression on same prediction data using Tensorflow